

## 8. Linear models for regression

Now we consider the problem of learning a function where the output is continuous:

$$f: X \rightarrow Y$$

$$X \subseteq \mathbb{R}^d$$

$$Y \subseteq \mathbb{R} \text{ + real numbers}$$

$$\text{train data set } \Delta = \left\{ (x_n, t_n)_{n=1}^N \right\}$$

Our dataset is still a set of pairs,  $t_n$  now can be a real value


The approximation of this function should return a real value. The easy way to define a real model, we define a linear combination of weights with the component vector input:

$$Y(\vec{x}; \vec{w}) = w_0 + w_1 x_1 + \dots + w_d x_d = \vec{w}^T \vec{x}$$

$$\vec{x}^T = \begin{pmatrix} 1 & x_1 & \dots & x_d \end{pmatrix}$$


dummy component to include  $w_0$  (bias component)

$$\vec{w}^T = (w_0 \ w_1 \ \dots \ w_d)$$



2-Dimensional Classification Problem

$$f: \mathbb{R}^2 \rightarrow \{+, o\}$$



1-Dimensional Regression problem representation

$$f: \mathbb{R} \rightarrow \mathbb{R}$$

There is a difference

With linear combinations we can learn only linear function. We can use the linear bias function model to generate LINEAR MODELS of NON-LINEAR - FUNCTIONS.

$$y(\vec{x}; \vec{w}) = \sum_{j=0}^M w_j \phi_j(\vec{x}) = \vec{w}^T \Phi(\vec{x})$$

with  $\vec{w}^T = (w_0 \dots w_M)^T$

$$\Phi(\vec{x}) = \begin{pmatrix} \phi_0(\vec{x}) \\ \vdots \\ \phi_M(\vec{x}) \end{pmatrix}$$

$\rightarrow 1$

FROM  $\vec{w}^0$  - point of view is a LINEAR MODEL

Since we are interested in computing  $\vec{w}^0$  we can apply the same reasoning of linear models

We consider a transformation of the input values,

is a linear combination of the weights and the transf. of the input by using the function  $\Phi$  (non-linear)

The fact that the function is now linear in  $\vec{x}$  does not affect the solution. As an example we can use as BASIS function a polynomial:

$$y = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$

# Linear Models for Regression

## Linear Basis Function Models

Using nonlinear functions of input variables:

$$y(\mathbf{x}; \mathbf{w}) = \sum_{j=0}^M w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}),$$



$$\text{with } \mathbf{w} = \begin{bmatrix} w_0 \\ \vdots \\ w_M \end{bmatrix}, \boldsymbol{\phi}(\mathbf{x}) = \begin{bmatrix} \phi_0(\mathbf{x}) \\ \vdots \\ \phi_M(\mathbf{x}) \end{bmatrix}, \text{ and } \phi_0(\mathbf{x}) = 1.$$

- Still linear in the parameters  $\mathbf{w}$ !

## Example: Polynomial curve fitting



$$y = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$

only  
 $w_0$



A line with  
 $y = w_0 + w_1 x$

better,  
quite good  
approx.



Very special case  
because the points  
are 10, you have  
a function that  
makes 0 error

**Warning: overfitting!!!**

If we just consider the error (of course we don't know the real function), the function with  $M=9$  is the best, but we don't like it because of OVERFITTING. We prefer the function with  $M=3$  (case of UNDERFITTING, we can see the shape of the data).

We can use different basis functions. How can compute parameters of our model? We will maximize the error.

The target value  $t$  is given by  $y(\vec{x}; \vec{w})$

Affected by additive noise  $\varepsilon$ :

$$t = y(\vec{x}; \vec{w}) + \varepsilon$$

We define an error model, assuming samples in dataset are not perfect

We assume that the error  $\varepsilon$  is gaussian, with  $\mu=0$ .

The error is a white noise process:

$$p(\varepsilon | \beta) = \mathcal{N}(\varepsilon | 0, \beta^{-1}) \quad \text{with precision } \beta \text{ that is the inverse of the variance.}$$

We have:


$$p(t | \vec{x}, \vec{w}, \beta) = \mathcal{N}(t | y(\vec{x}; \vec{w}), \beta^{-1})$$

likelihood expressed with the Gaussian


exploited the fact that the Gaussian is additive

# Linear Regression Basis Functions


## Examples of basis functions



Polynomial



Radial



Sigmoid / Tanh

# Linear Regression - Algorithms

## Maximum likelihood and least squares

Target value  $t$  is given by  $y(\mathbf{x}; \mathbf{w})$  affected by additive noise  $\epsilon$

$$t = y(\mathbf{x}; \mathbf{w}) + \epsilon$$

Assume Gaussian noise  $P(\epsilon|\beta) = \mathcal{N}(\epsilon|0, \beta^{-1})$ , with precision (inverse variance)  $\beta$ .

We have:

$$P(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t|y(\mathbf{x}; \mathbf{w}), \beta^{-1})$$

We want to compute  $\vec{w}$  given the dataset.

Assuming the observations are independent and identically distributed, we can compute the probability in this way:

$$p(\{t_1, \dots, t_N\} \mid x_1, \dots, x_N, \vec{w}, \beta) =$$
$$= \prod_{n=1}^N \frac{1}{\sqrt{\beta}} \exp\left(-\frac{1}{2\beta} (\vec{w}^T \phi(x_n) - t_n)^2\right)$$

Changing  $\vec{w}$  you change the curve.

And we want the value of  $\vec{w}$  for which the probability distribution is maximum. We solve the problem again by defining an error function:

$$\ln(p(\{t_1, \dots, t_N\} \mid x_1, \dots, x_N, \vec{w}, \beta)) =$$
$$= \sum \ln \frac{1}{\sqrt{\beta}} \exp\left(-\frac{1}{2\beta} (\vec{w}^T \phi(x_n) - t_n)^2\right) =$$
$$= -\frac{1}{2} \sum_{n=1}^N (\vec{w}^T \phi(x_n) - t_n)^2 - \frac{N}{2} \ln(2\pi\beta)$$

$E_D(\vec{w})$

↳ error function: it has negative sign

Maximum likelihood  $\Rightarrow$  minimization of  $E_D(\vec{w})$

In this case it corresponds to least square error minimization, but  $t_n$  here are real values

## Linear Regression - Algorithms

Assume observations independent and identically distributed (i.i.d.)

We seek the maximum of the likelihood function:

$$P(\{t_1, \dots, t_N\} | \mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}).$$

or equivalently:

$$\begin{aligned} \ln P(\{t_1, \dots, t_N\} | \mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{w}, \beta) &= \sum_{n=1}^N \ln \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) \\ &= -\beta \underbrace{\frac{1}{2} \sum_{n=1}^N [t_n - \mathbf{w}^T \phi(\mathbf{x}_n)]^2}_{E_D(\mathbf{w})} - \frac{N}{2} \ln(2\pi\beta^{-1}). \end{aligned}$$

## Linear Regression - Algorithms

Maximum likelihood (zero-mean Gaussian noise assumption)

$$\operatorname{argmax} P(\{t_1, \dots, t_N\} | \mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{w}, \beta)$$

corresponds to least square error minimization

$$\operatorname{argmin} E_D(\mathbf{w}) = \operatorname{argmin} \frac{1}{2} \sum_{n=1}^N [t_n - \mathbf{w}^T \phi(\mathbf{x}_n)]^2$$

# Linear Regression - Algorithms

Note:

$$E_D(\mathbf{w}) = \frac{1}{2}(\mathbf{t} - \Phi\mathbf{w})^T(\mathbf{t} - \Phi\mathbf{w}),$$

$$\text{with } \mathbf{t} = \begin{bmatrix} t_1 \\ \vdots \\ t_N \end{bmatrix} \text{ and } \Phi = \begin{bmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \cdots & \phi_M(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \cdots & \phi_M(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \cdots & \phi_M(\mathbf{x}_N) \end{bmatrix}.$$

Optimality condition:

$$\nabla E_D = 0 \iff \Phi^T \Phi \mathbf{w} = \Phi^T \mathbf{t}.$$

Hence:

$$\mathbf{w}_{ML} = \underbrace{(\Phi^T \Phi)^{-1} \Phi^T}_{\Phi^\dagger: \text{pseudo-inverse}} \mathbf{t}.$$

# Linear Regression - Algorithms

## Sequential Learning

Stochastic gradient descent algorithm:

$$\hat{\mathbf{w}} \leftarrow \hat{\mathbf{w}} - \eta \nabla E_n$$

$\eta$ : learning rate parameter

Therefore:

$$\hat{\mathbf{w}} \leftarrow \hat{\mathbf{w}} + \eta [t_n - \hat{\mathbf{w}}^T \phi(\mathbf{x}_n)] \phi(\mathbf{x}_n)$$

Algorithm converges for suitable small values of  $\eta$ .



# Linear Regression - Regularization

Regularization is a technique to control over-fitting.

$$\operatorname{argmin} E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

with  $\lambda > 0$  being the regularization factor

A common choice:


$$E_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}.$$

Other choices:

$$E_W(\mathbf{w}) = \sum_{j=0}^M |w_j|^q.$$



# Linear Regression - Regularization

$$\operatorname{argmin} E_D(\mathbf{w})$$



# Linear Regression - Regularization

$$\operatorname{argmin} E_D(\mathbf{w}) + \lambda \frac{1}{2} \mathbf{w}^T \mathbf{w}$$



## Linear Regression - Multiple outputs

$\mathbf{y}$ : vector with  $K$  components

$$\mathbf{y}(\mathbf{x}; \mathbf{W}) = \mathbf{W}^T \phi(\mathbf{x})$$

Target variable  $\mathbf{T}$ , with  $\mathbf{t}_n$  vector of  $K$  output values for input  $\mathbf{x}_n$

$$\ln P(\mathbf{T}|\mathbf{X}, \mathbf{W}, \beta) = \sum_{n=1}^N \ln \mathcal{N}(\mathbf{t}_n | \mathbf{W}^T \phi(\mathbf{x}_n), \beta^{-1} \mathbf{I})$$

Similarly as before we obtain:

$$\mathbf{W}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{T}.$$